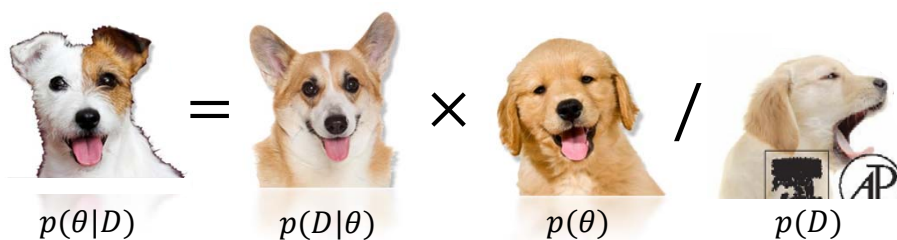


Doing Bayesian Data Analysis


$$p(\theta|D) = p(D|\theta) \times p(\theta) / p(D)$$

John K. Kruschke

© John K. Kruschke, Oct. 2012

1

Outline of Talk:

- Bayesian reasoning generally.
- Bayesian estimation applied to two groups. Rich information.
- The NHST t test: perfidious p values and the con game of confidence intervals.
- Conclusion: Bayesian estimation supersedes NHST.

© John K. Kruschke, Oct. 2012

2

Bayesian Reasoning

The role of data is to re-allocate credibility:

Prior Credibility with **New Data**
→ **Posterior Credibility**

via Bayes' rule

© John K. Kruschke, Oct. 2012

3

Bayesian Reasoning

The role of data is to re-allocate credibility:

**Bayesian reasoning in everyday life is
intuitive:**

© John K. Kruschke, Oct. 2012

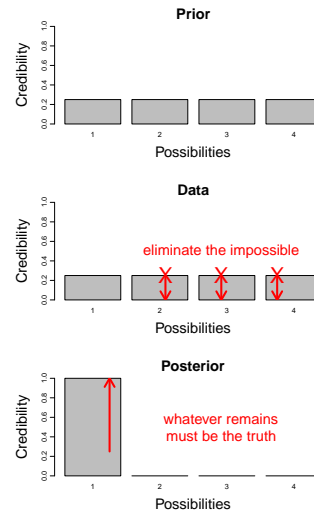
4

Bayesian Reasoning

The role of data is to re-allocate credibility:

Bayesian reasoning in everyday life is intuitive:

Sherlock Holmes: "How often have I said to you that when you have eliminated the impossible, whatever remains, however improbable, must be the truth?" (Doyle, 1890)



© John K. Kruschke, Oct. 2012

5

Bayesian Reasoning

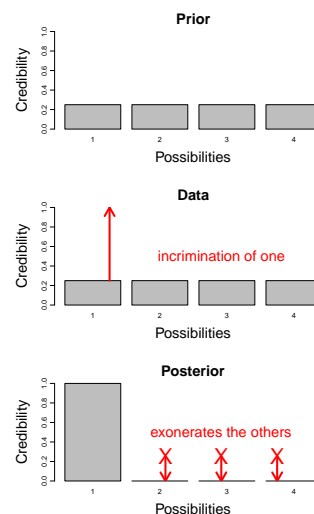
The role of data is to re-allocate credibility:

Bayesian reasoning in everyday life is intuitive:

Sherlock Holmes: "How often have I said to you that when you have eliminated the impossible, whatever remains, however improbable, must be the truth?" (Doyle, 1890)

Judicial exoneration: For unaffiliated suspects, the incrimination of one exonerates the others.

Credibility of the claim that the suspect committed the crime.



© John K. Kruschke, Oct. 2012

6

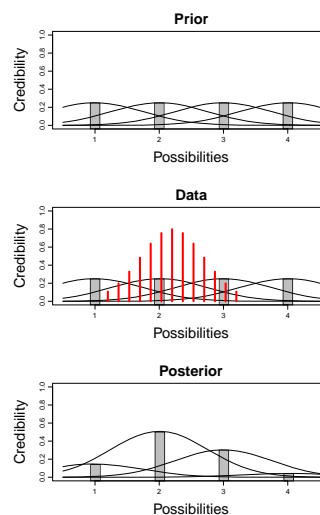
Bayesian Data Analysis

The role of data is to re-allocate credibility:

Bayesian reasoning in data analysis is intuitive:

Possibilities are *parameter values* in a model, such as the *mean* of a normal distribution.

We reallocate credibility to parameter values that are consistent with the data.



© John K. Kruschke, Oct. 2012

7

Bayesian Data Analysis

The role of data is to re-allocate credibility:

1. Define a meaningful descriptive model.
2. Establish prior credibility regarding parameter values in the model. The prior credibility must be acceptable to a skeptical scientific audience.
3. Collect data.
4. Use Bayes' rule to re-allocate credibility to parameter values that are most consistent with the data.

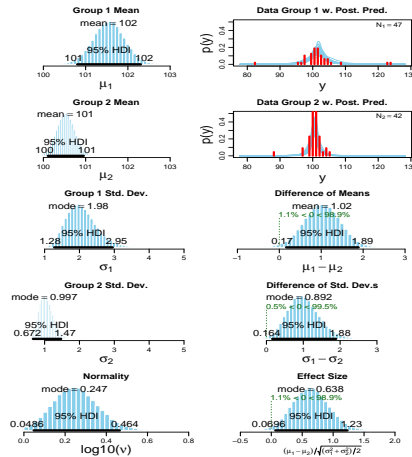
© John K. Kruschke, Oct. 2012

8

Robust Bayesian estimation for comparing two groups

Consider two groups;
e.g.,
IQ of “smart drug” group
and of control group.

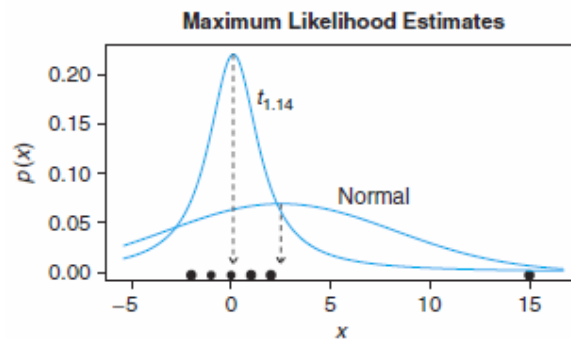
Step 1: Define a model
for describing the data.



© John K. Kruschke, Oct. 2012

10

Descriptive distribution for data with outliers




Normal is
pulled by
outliers, but t
distribution is
not.

t distribution is used here as a description of data,
NOT as a sampling distribution for p values!

© John K. Kruschke, Oct. 2012

11

Descriptive distribution for data with outliers




The t distribution has normality controlled by the parameter v .

© John K. Kruschke, Oct. 2012

13

Robust Bayesian estimation for comparing two groups




The data from each group are described by t distributions, using five parameters altogether.

© John K. Kruschke, Oct. 2012

14


Robust Bayesian estimation for comparing two groups



© John K. Kruschke, Oct. 2012

15


Robust Bayesian estimation for comparing two groups



© John K. Kruschke, Oct. 2012

16


Robust Bayesian estimation for comparing two groups



© John K. Kruschke, Oct. 2012

17


Robust Bayesian estimation for comparing two groups



© John K. Kruschke, Oct. 2012

18

Robust Bayesian estimation for comparing two groups




Parameter distributions will be represented by histograms: A huge number of representative parameter values.

© John K. Kruschke, Oct. 2012

19

Step 3: Collect Data.

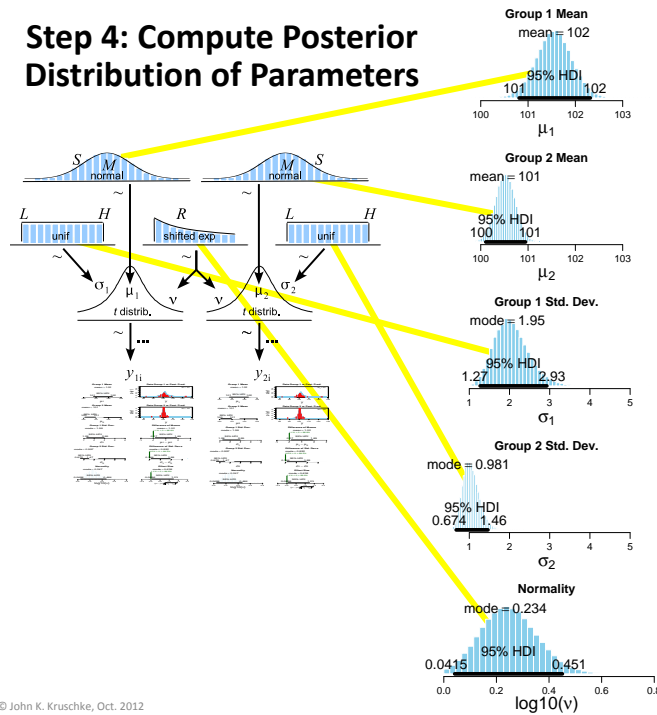


One fixed data set,
shown as red
histograms.

© John K. Kruschke, Oct. 2012

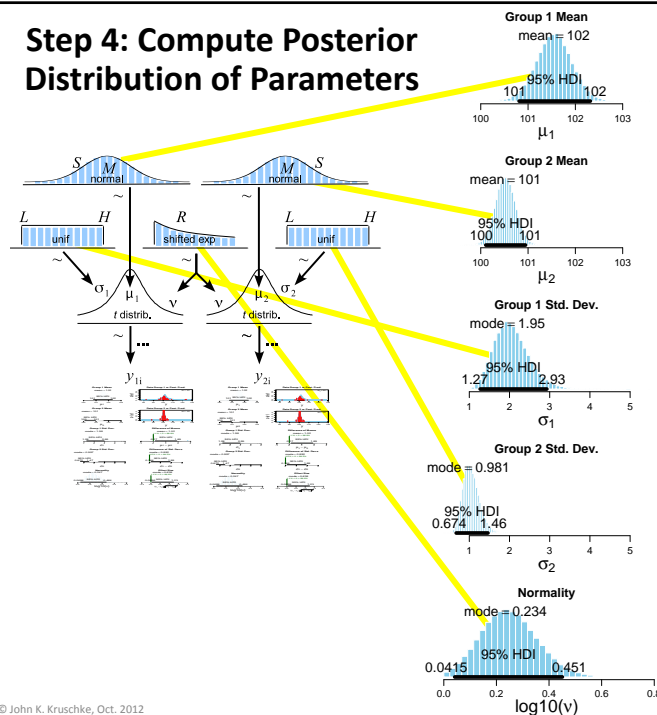
20

Step 4: Compute Posterior Distribution of Parameters



© John K. Kruschke, Oct. 2012

Step 4: Compute Posterior Distribution of Parameters



© John K. Kruschke, Oct. 2012

Important:
These are histograms of parameter values from the posterior distribution:
A huge number of combinations of $\mu_1, \mu_2, \sigma_1, \sigma_2, v$ that are jointly credible given the data.

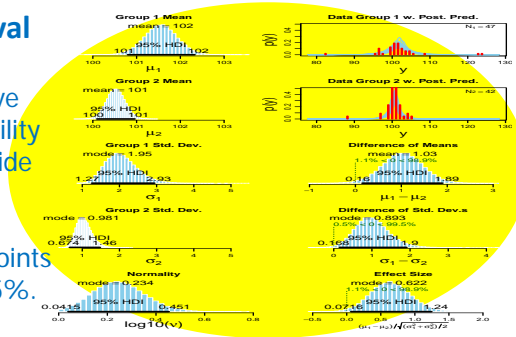
These are *not* data distributions, and *not* sampling distributions from a null hypothesis.

95% HDI: Highest density interval

Points within the HDI have higher credibility (probability density) than points outside the HDI.

The total probability of points within the 95% HDI is 95%.

Points outside the HDI may be deemed not credible.



© John K. Kruschke, Oct. 2012

23

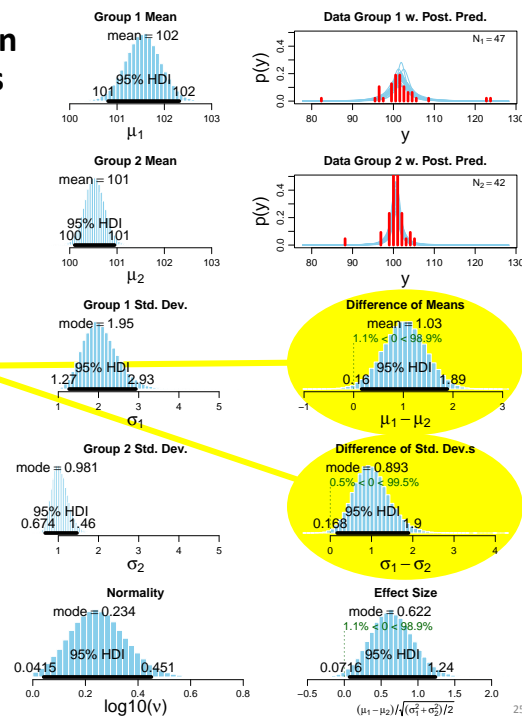
Robust Bayesian estimation for comparing two groups

Differences between
groups?

Compute $\mu_1 - \mu_2$
and $\sigma_1 - \sigma_2$
at each of the many
credible combinations.

Here, both differences
are credibly non-zero.

(NHST would require
two tests...)



© John K. Kruschke, Oct. 2012

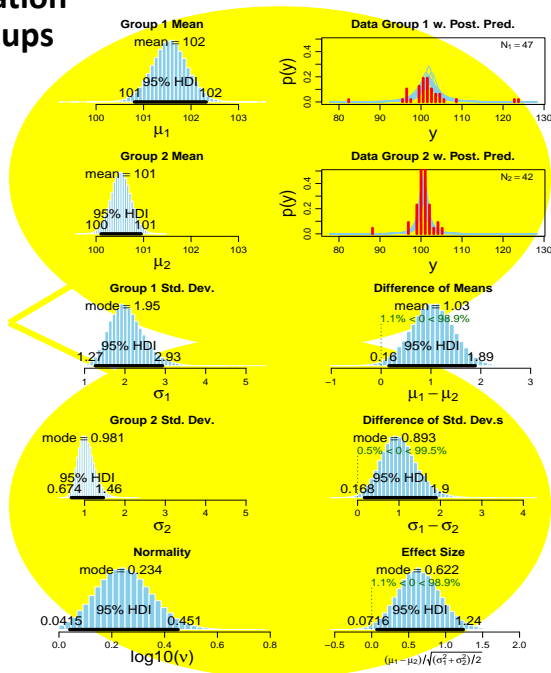
25

Robust Bayesian estimation for comparing two groups

Differences between
groups?
Compute $\mu_1 - \mu_2$
and $\sigma_1 - \sigma_2$
at each of the many
credible combinations.

Here, both differences
are credibly non-zero.

(NHST would require
two tests...)

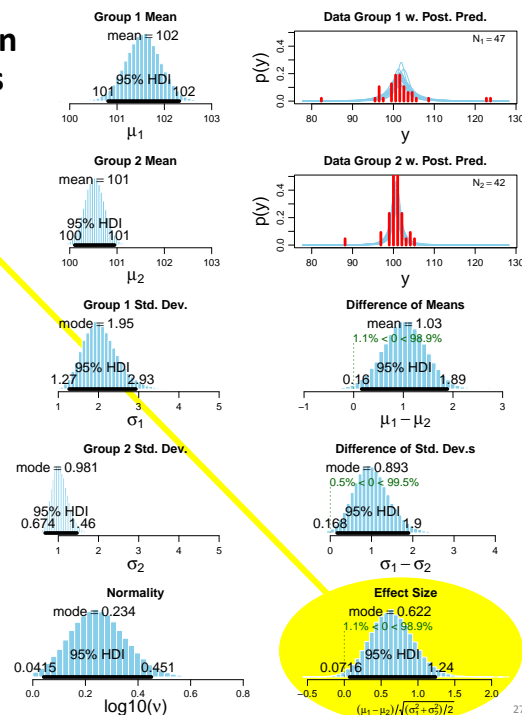


© John K. Kruschke, Oct. 2012

26

Robust Bayesian estimation for comparing two groups

Complete distribution
on effect size!




© John K. Kruschke, Oct. 2012

27

Robust Bayesian estimation for comparing two groups

Complete distribution
on effect size!



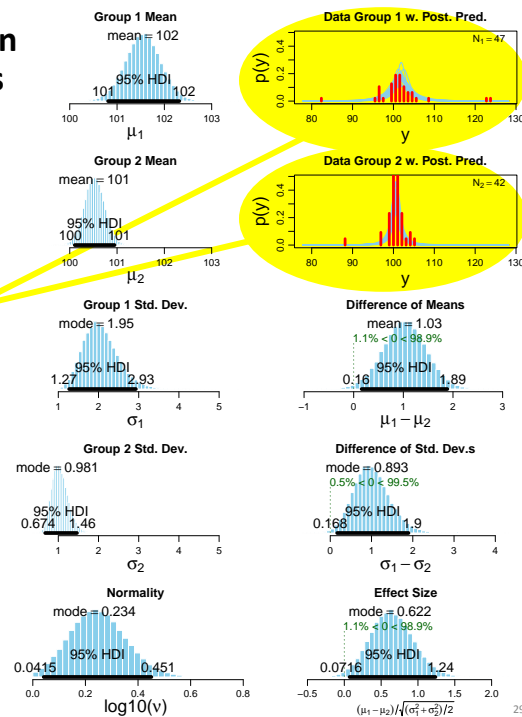
© John K. Kruschke, Oct. 2012

28

Robust Bayesian estimation for comparing two groups

Are the data described
well by the model?

Superimpose a
smattering of credible
descriptive distributions
on data.
= “posterior predictive
check”



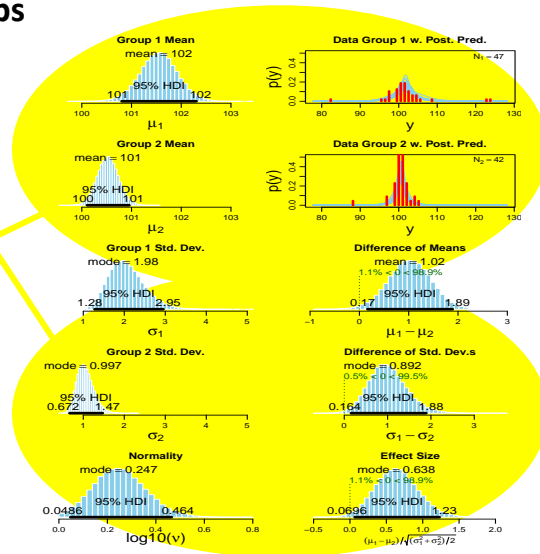
© John K. Kruschke, Oct. 2012

29

Robust Bayesian estimation for comparing two groups

Are the data described well by the model?

Superimpose a smattering of credible descriptive distributions on data.
= “posterior predictive check”



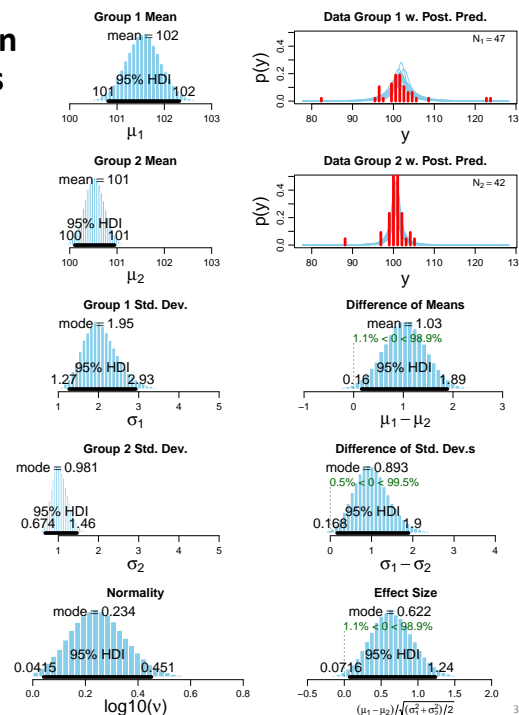
© John K. Kruschke, Oct. 2012

30

Robust Bayesian estimation for comparing two groups

Summary:

- Complete distribution of credible parameter values (not merely point estimate with ends of confidence interval).
- Decisions about multiple aspects of parameters (without reference to p values).
- Flexible descriptive model, robust to outliers (unlike NHST t test).



© John K. Kruschke, Oct. 2012

31

Computer Software:

Packaged for easy use!
Underlying program is never seen.

```
source("BEST.R") # load the program

# Specify data as vectors (replace with your own data):
y1 = c(101,100,102,104,102,97,105,105,98,101,100,123,105,
       109,102,82,102,100,102,102,101,102,102,103,103,97,
       96,103,124,101,101,100,101,101,104,100,101)
y2 = c(99,101,100,101,102,100,97,101,104,101,102,102,100,
       104,100,100,100,101,102,103,97,101,101,100,101,99,
       101,100,99,101,100,102,99,100,99)

# Run the Bayesian analysis:
mcmcChain = BESTmcmc( y1 , y2 )

# Plot the results of the Bayesian analysis:
BESTplot( y1 , y2 , mcmcChain )
```


© John K. Kruschke, Oct. 2012

32

Robust Bayesian estimation for comparing two groups

Download the programs from
<http://www.indiana.edu/~kruschke/BEST/BEST.zip>

Now for a look
under the hood



<http://www.autonationconnect.com/2010/07/backseat-mechanic-under-the-hood/>

© John K. Kruschke, Sept. 2012

33

Doing it with JAGS

“JAGS” = Just Another Gibbs Sampler
but other sampling methods are incorporated.

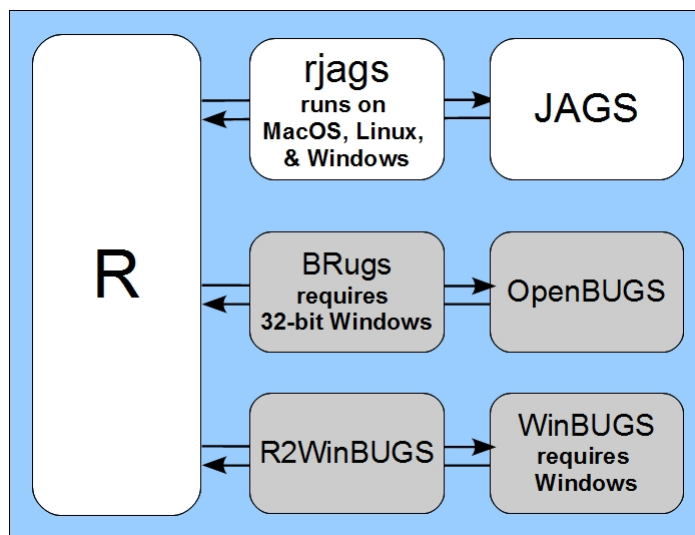


JAGS makes it easy. You specify only the

- prior function
- likelihood function

and JAGS does the rest! You do no math, no selection of sampling methods.

JAGS and BUGS



Installation: See Blog Entry

<http://doingbayesiandataanalysis.blogspot.com/2012/01/complete-steps-for-installing-software.html>

Doing Bayesian Data Analysis

JAGS! Bayes is Better Software Tips Workshops Book Reviews Bayesian Jokes Suggestion Box

Saturday, January 28, 2012

Complete steps for installing software and programs

To use the programs, there are five basic installation steps:

1. Install the general programming language R. Go to the [R web site](#) and install the latest version of R appropriate for your computer (Windows, Mac, Linux).
2. Install the Bayesian sampling program JAGS. Go to the [JAGS web site](#) and install the latest version of JAGS appropriate for your computer (Windows, Mac, Linux). Caution to Mac users: If you find that R is having trouble working with JAGS, it might be because x2app (in the next step) is compiled for use with JAGS 2.*. Instead of JAGS 3.*, if you have this problem, uninstall JAGS 3.* and install JAGS 2.* instead.
3. Install the package that lets R talk to JAGS. Invoke R. At the command line, type `install.packages("rjags")`. You may be prompted to select an Internet archive to get the package from; select a site geographically near you.
4. Get the programs used in the book. For a list of individually downloadable programs, click [HERE](#). For a zip file (named `ProgramsDoingBayesianDataAnalysis.zip`) that contains all the programs in a single file, click [HERE](#). Be sure to unzip (extract) the programs from the zip file. Programs are updated occasionally. Click on the list link, above, and sort the list by date modified to see if you have the most recent versions. I now recommend using JAGS instead of BUGS. All the programs now have JAGS versions in addition to the original BUGS versions. Any program that had "bugs" or "brugs" in its file name now has a JAGS equivalent with "jags" in its file name. More info is available at [this blog post](#).
5. Install the R editor RStudio. It comes with its own built-in editor, but it is not very useful for dealing with long programs. Go to the [RStudio web site](#) and install the latest version of RStudio appropriate for your computer (Windows, Mac, Linux).

Tips for running the programs:

- Many of the programs call others when running, so
- put all the programs together in the same folder and
- be sure that it has that folder as its working directory. Set the working directory in RStudio from the menu: Tools -> Set Working Dir. Or, if you're using R's built-in editor, set the working directory in R's command console from the menu: File -> Change dir...

Mac users: Having trouble with graphics when running the programs? See [this blog post](#).

The Book (click it for info):

Doing Bayesian Data Analysis: A Tutorial with R and BUGS

Search This Blog

Blog Archive

2012 (15)

February (5)

January (10)

Complete steps for installing software and program...

Jocular Disbelief


Specialized Workshop, Feb 17-18, Seton Hall Univer...

© John K. Kruschke, Sept. 2012

Robust Bayesian estimation for comparing two groups

Program BEST.R: JAGS model specification.


```
model {
  for ( i in 1:Ntotal ) {
    y[i] ~ dt( mu[x[i]], tau[x[i]] , nu )
  }
  for ( j in 1:2 ) {
    mu[j] ~ dnorm( muM , muP )
    tau[j] <- 1/pow( sigma[j] , 2 )
    sigma[j] ~ dunif( sigmaLow , sigmaHigh )
  }
  nu <- nuMinusOne+1
  nuMinusOne ~ dexp(1/29)
}
```



Robust Bayesian estimation for comparing two groups

Program BEST.R: JAGS model specification.

```
model {
  for ( i in 1:Ntotal ) {
    y[i] ~ dt( mu[x[i]], tau[x[i]], nu )
  }
  for ( j in 1:2 ) {
    mu[j] ~ dnorm( muM , muP )
    tau[j] <- 1/pow( sigma[j] , 2 )
    sigma[j] ~ dunif( sigmaLow , sigmaHigh )
  }
  nu <- nuMinusOne+1
  nuMinusOne ~ dexp(1/29)
}
```




© John K. Kruschke, Sept. 2012

38

Robust Bayesian estimation for comparing two groups

Program BEST.R: JAGS model specification.

```
model {
  for ( i in 1:Ntotal ) {
    y[i] ~ dt( mu[x[i]], tau[x[i]], nu )
  }
  for ( j in 1:2 ) {
    mu[j] ~ dnorm( muM , muP )
    tau[j] <- 1/pow( sigma[j] , 2 )
    sigma[j] ~ dunif( sigmaLow , sigmaHigh )
  }
  nu <- nuMinusOne+1
  nuMinusOne ~ dexp(1/29)
}
```



Nested indexing:
x[i] is the group (1 or 2)
of the ith score.


© John K. Kruschke, Sept. 2012

39

Robust Bayesian estimation for comparing two groups

Program BEST.R: JAGS model specification.

```
model {
  for ( i in 1:Ntotal ) {
    y[i] ~ dt( mu[x[i]], tau[x[i]], nu )
  }
  for ( j in 1:2 ) {
    mu[j] ~ dnorm( muM , muP )
    tau[j] <- 1/pow( sigma[j] , 2 )
    sigma[j] ~ dunif( sigmaLow , sigmaHigh )
  }
  nu <- nuMinusOne+1
  nuMinusOne ~ dexp(1/29)
}
```




© John K. Kruschke, Sept. 2012

40

Robust Bayesian estimation for comparing two groups

Program BEST.R: JAGS model specification.

```
model {
  for ( i in 1:Ntotal ) {
    y[i] ~ dt( mu[x[i]], tau[x[i]], nu )
  }
  for ( j in 1:2 ) {
    mu[j] ~ dnorm( muM , muP )
    tau[j] <- 1/pow( sigma[j] , 2 )
    sigma[j] ~ dunif( sigmaLow , sigmaHigh )
  }
  nu <- nuMinusOne+1
  nuMinusOne ~ dexp(1/29)
}
```




© John K. Kruschke, Sept. 2012

41

Robust Bayesian estimation for comparing two groups

Program BEST.R: JAGS model specification.

```
model {
  for ( i in 1:Ntotal ) {
    y[i] ~ dt( mu[x[i]], tau[x[i]], nu )
  }
  for ( j in 1:2 ) {
    mu[j] ~ dnorm( muM , muP )
    tau[j] <- 1/pow( sigma[j] , 2 )
    sigma[j] ~ dunif( sigmaLow , sigmaHigh )
  }
  nu <- nuMinusOne+1
  nuMinusOne ~ dexp(1/29)
}
```



© John K. Kruschke, Sept. 2012

42

Programs in R + rjags + JAGS:

Five main sections in all programs:

1. **Specify model** (we just did this).
2. **Load data.**
3. **Initialize the MCMC chain.**
4. **Run the MCMC chain.**
5. **Examine the results.**

© John K. Kruschke, Sept. 2012

43

BEST.R

```

BESTmcmc = function( y1, y2, numSavedSteps=100000, thinSteps=1, showMCMC=FALSE) {
  # This function generates an MCMC sample from the posterior distribution.
  # Description of arguments:
  # showMCMC is a flag for displaying diagnostic graphs of the chains.
  # If F (the default), no chain graphs are displayed. If T, they are.

  require(rjags)

  #-----
  # THE MODEL.
  modelString = "
  model {
    for ( i in 1:Ntotal ) {
      y[i] ~ dt( mu[x[i]], tau[x[i]] , nu )
    }
    for ( j in 1:2 ) {
      mu[j] ~ dnorm( muM , muP )
      tau[j] <- 1/pow( sigma[j] , 2 )
      sigma[j] ~ dunif( sigmaLow , sigmaHigh )
    }
    nu <- nuMinusOne+1
    nuMinusOne ~ dexp(1/29)
  }
  " # close quote for modelString
  # Write out modelString to a text file
  writeLines( modelString , con="BESTmodel.txt" )

  #-----
  # THE DATA.
  # Load the data:
  y = c( y1 , y2 ) # combine data into one vector
  x = c( rep(1,length(y1)) , rep(2,length(y2)) ) # create group membership code
  Ntotal = length(y)
  # Specify the data in a list, for later shipment to JAGS:
  dataList = list(
    y = y ,
    x = x ,
    Ntotal = Ntotal ,

```

© John K. Kruschke, Sept. 2012

44

BEST.R

```

BESTmcmc = function( y1, y2, numSavedSteps=100000, thinSteps=1, showMCMC=FALSE) {
  # This function generates an MCMC sample from the posterior distribution.
  # Description of arguments:
  # showMCMC is a flag for displaying diagnostic graphs of the chains.
  # If F (the default), no chain graphs are displayed. If T, they are.

  require(rjags)

  #-----
  # THE MODEL.
  modelString = "
  model {
    for ( i in 1:Ntotal ) {
      y[i] ~ dt( mu[x[i]], tau[x[i]] , nu )
    }
    for ( j in 1:2 ) {
      mu[j] ~ dnorm( muM , muP )
      tau[j] <- 1/pow( sigma[j] , 2 )
      sigma[j] ~ dunif( sigmaLow , sigmaHigh )
    }
    nu <- nuMinusOne+1
    nuMinusOne ~ dexp(1/29)
  }
  " # close quote for modelString
  # Write out modelString to a text file
  writeLines( modelString , con="BESTmodel.txt" )

  #-----
  # THE DATA.
  # Load the data:
  y = c( y1 , y2 ) # combine data into one vector
  x = c( rep(1,length(y1)) , rep(2,length(y2)) ) # create group membership code
  Ntotal = length(y)
  # Specify the data in a list, for later shipment to JAGS:
  dataList = list(
    y = y ,
    x = x ,
    Ntotal = Ntotal ,

```

© John K. Kruschke, Sept. 2012

45

```

    }
    nu <- nuMinusOne+1
    nuMinusOne ~ dexp(1/29)
  }
  " # close quote for modelString
  # Write out modelString to a text file
  writeLines( modelString , con="BESTmodel.txt" )

  #-----
  # THE DATA.
  # Load the data:
  y = c( y1 , y2 ) # combine data into one vector
  x = c( rep(1,length(y1)) , rep(2,length(y2)) ) # create group membership code
  Ntotal = length(y)
  # Specify the data in a list, for later shipment to JAGS:
  dataList = list(
    y = y ,
    x = x ,
    Ntotal = Ntotal ,
    muM = mean(y) ,
    muP = 0.000001 * 1/sd(y)^2 ,
    sigmaLow = sd(y) / 1000 ,
    sigmaHigh = sd(y) * 1000
  )

  #-----
  # INITIALIZE THE CHAINS.
  # Initial values of MCMC chains based on data:
  mu = c( mean(y1) , mean(y2) )
  sigma = c( sd(y1) , sd(y2) )
  # Regarding initial values in next line: (1) sigma will tend to be too big if
  # the data have outliers, and (2) nu starts at 5 as a moderate value. These
  # initial values keep the burn-in period moderate.
  initsList = list( mu = mu , sigma = sigma , nuMinusOne = 4 )

  #-----
  # RUN THE CHAINS

  parameters = c( "mu" , "sigma" , "nu" ) # The parameters to be monitored
  adaptSteps = 500 # Number of steps to "tune" the samplers
  burnInSteps = 1000
  nChains = 3
  # After = ceiling( ( sumSavedSteps + thinSteps ) / nChains )

```

© John K. Kruschke, Sept. 2012 46

```

    }
    nu <- nuMinusOne+1
    nuMinusOne ~ dexp(1/29)
  }
  " # close quote for modelString
  # Write out modelString to a text file
  writeLines( modelString , con="BESTmodel.txt" )

  #-----
  # THE DATA.
  # Load the data:
  y = c( y1 , y2 ) # combine data into one vector
  x = c( rep(1,length(y1)) , rep(2,length(y2)) ) # create group membership code
  Ntotal = length(y)
  # Specify the data in a list, for later shipment to JAGS:
  dataList = list(
    y = y ,
    x = x ,
    Ntotal = Ntotal ,
    muM = mean(y) ,
    muP = 0.000001 * 1/sd(y)^2 ,
    sigmaLow = sd(y) / 1000 ,
    sigmaHigh = sd(y) * 1000
  )

  #-----
  # INITIALIZE THE CHAINS.
  # Initial values of MCMC chains based on data:
  mu = c( mean(y1) , mean(y2) )
  sigma = c( sd(y1) , sd(y2) )
  # Regarding initial values in next line: (1) sigma will tend to be too big if
  # the data have outliers, and (2) nu starts at 5 as a moderate value. These
  # initial values keep the burn-in period moderate.
  initsList = list( mu = mu , sigma = sigma , nuMinusOne = 4 )

  #-----
  # RUN THE CHAINS

  parameters = c( "mu" , "sigma" , "nu" ) # The parameters to be monitored
  adaptSteps = 500 # Number of steps to "tune" the samplers
  burnInSteps = 1000
  nChains = 3
  # After = ceiling( ( sumSavedSteps + thinSteps ) / nChains )

```

© John K. Kruschke, Sept. 2012 47

```

# initial values keep the burn-in period moderate.
initsList = list( mu = mu , sigma = sigma , nuMinusOne = 4 )

#-----
# RUN THE CHAINS

parameters = c( "mu" , "sigma" , "nu" )      # The parameters to be monitored
adaptSteps = 500                             # Number of steps to "tune" the samplers
burnInSteps = 1000
nChains = 3
nIter = ceiling( ( numSavedSteps * thinSteps ) / nChains )
# Create, initialize, and adapt the model:
jagsModel = jags.model( "BESTmodel.txt" , data=dataList , inits=initsList ,
                        n.chains=nChains , n.adapt=adaptSteps )

# Burn-in:
cat( "Burning in the MCMC chain...\n" )
update( jagsModel , n.iter=burnInSteps )
# The saved MCMC chain:
cat( "Sampling final MCMC chain...\n" )
codaSamples = coda.samples( jagsModel , variable.names=parameters ,
                           n.iter=nIter , thin=thinSteps )

# resulting codaSamples object has these indices:
# codaSamples[ [ chainIdx ] ][ stepIdx , paramIdx ]

#-----
# EXAMINE THE RESULTS
if ( showMCMC ) {
  windows()
  autocorr.plot( codaSamples[[1]] , ask=FALSE )
}

# Convert coda-object codaSamples to matrix object for easier handling.
# But note that this concatenates the different chains into one long chain.
# Result is mcmcChain[ stepIdx , paramIdx ]
mcmcChain = as.matrix( codaSamples )
return( mcmcChain )

} # end function BESTmcmc

```

© John K. Kruschke, Sept. 2012

48

```

# initial values keep the burn-in period moderate.
initsList = list( mu = mu , sigma = sigma , nuMinusOne = 4 )

#-----
# RUN THE CHAINS

parameters = c( "mu" , "sigma" , "nu" )      # The parameters to be monitored
adaptSteps = 500                             # Number of steps to "tune" the samplers
burnInSteps = 1000
nChains = 3
nIter = ceiling( ( numSavedSteps * thinSteps ) / nChains )
# Create, initialize, and adapt the model:
jagsModel = jags.model( "BESTmodel.txt" , data=dataList , inits=initsList ,
                        n.chains=nChains , n.adapt=adaptSteps )

# Burn-in:
cat( "Burning in the MCMC chain...\n" )
update( jagsModel , n.iter=burnInSteps )
# The saved MCMC chain:
cat( "Sampling final MCMC chain...\n" )
codaSamples = coda.samples( jagsModel , variable.names=parameters ,
                           n.iter=nIter , thin=thinSteps )

# resulting codaSamples object has these indices:
# codaSamples[ [ chainIdx ] ][ stepIdx , paramIdx ]

#-----
# EXAMINE THE RESULTS
if ( showMCMC ) {
  windows()
  autocorr.plot( codaSamples[[1]] , ask=FALSE )
}

# Convert coda-object codaSamples to matrix object for easier handling.
# But note that this concatenates the different chains into one long chain.
# Result is mcmcChain[ stepIdx , paramIdx ]
mcmcChain = as.matrix( codaSamples )
return( mcmcChain )

} # end function BESTmcmc

```

© John K. Kruschke, Sept. 2012

49

Computer Software:

Packaged for easy use!
Underlying program is never seen.

```
source("BEST.R") # load the program

# Specify data as vectors (replace with your own data):
y1 = c(101,100,102,104,102,97,105,105,98,101,100,123,105,
       109,102,82,102,100,102,102,101,102,102,103,103,97,
       96,103,124,101,101,100,101,101,104,100,101)
y2 = c(99,101,100,101,102,100,97,101,104,101,102,102,100,
       104,100,100,100,101,102,103,97,101,101,100,101,99,
       101,100,99,101,100,102,99,100,99)

# Run the Bayesian analysis:
mcmcChain = BESTmcmc( y1 , y2 )

# Plot the results of the Bayesian analysis:
BESTplot( y1 , y2 , mcmcChain )
```

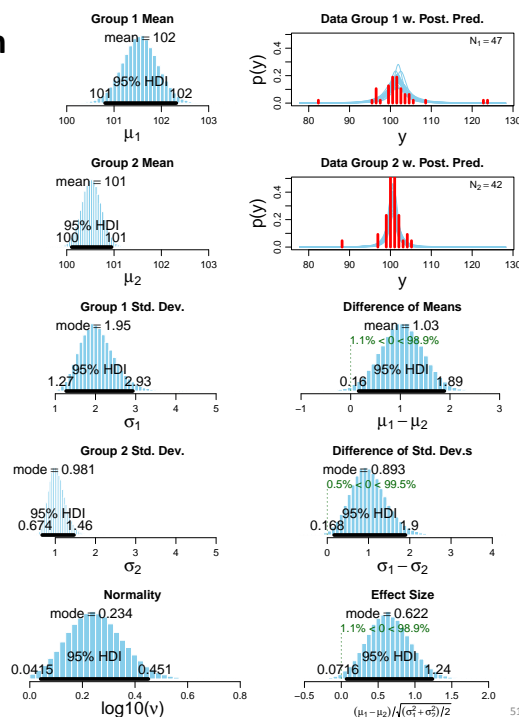
© John K. Kruschke, Sept. 2012

50

Recall Bayesian estimation for comparing two groups

Summary:

- Complete distribution of credible parameter values (not merely point estimate with ends of confidence interval).
- Decisions about multiple aspects of parameters (without reference to p values).
- Flexible descriptive model, robust to outliers (unlike NHST t test).



© John K. Kruschke, Oct. 2012

51

Recall Bayesian estimation:

What does NHST say?

t test of means:

$t(87)=1.62$, $p=0.110$ ($>.05$)


95%CI: -0.361 to 3.477.

F test of variances:

$F(46,41)=5.72$, $p < .001$.

95%CI on difference: ?

But, must apply corrections
for multiple tests.



Recall Bayesian estimation:

What does NHST say?

Oops! Data are not normal,
so do **resampling** instead.

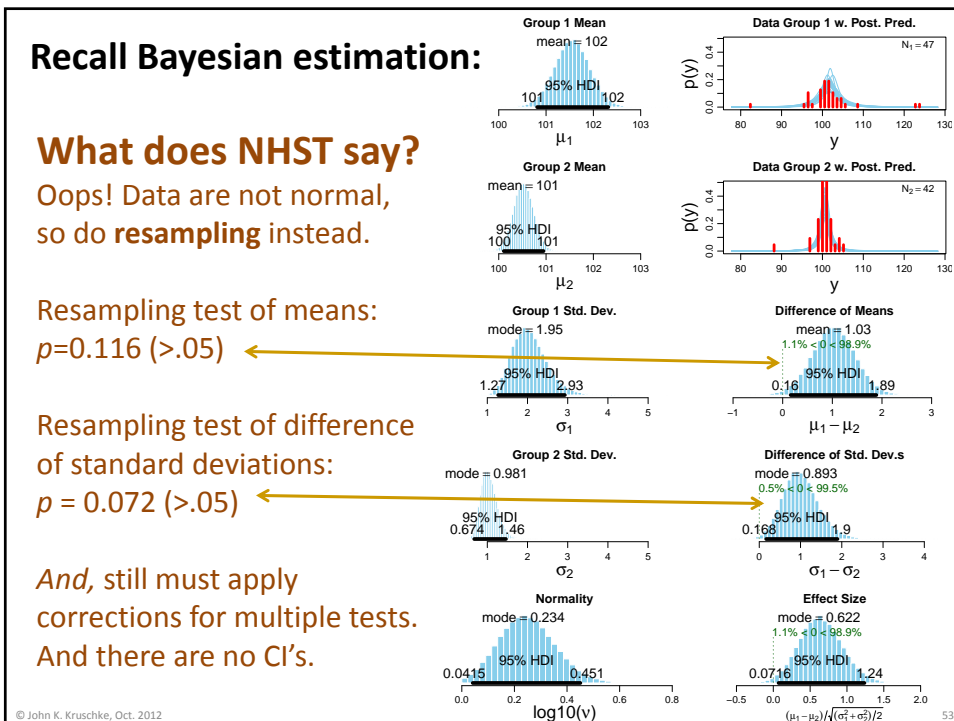
Resampling test of means:

$p=0.116$ ($>.05$)

Resampling test of difference
of standard deviations:

$p = 0.072$ ($>.05$)

And, still must apply
corrections for multiple tests.
And there are no CI's.



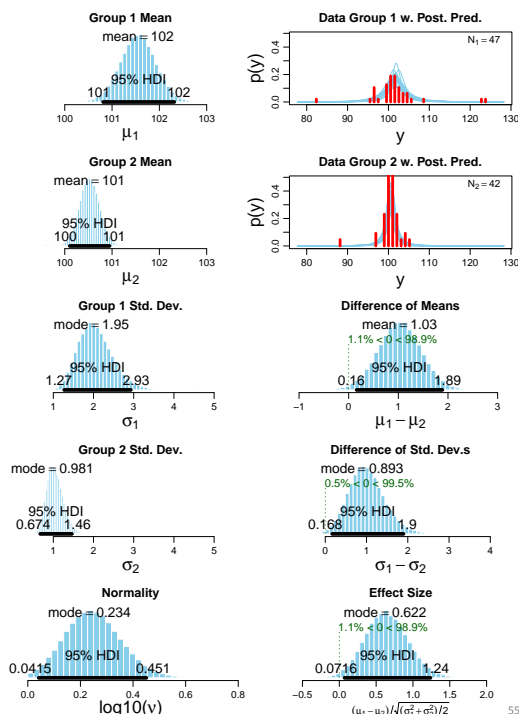
Example with outliers: BESTexample.R

Bayesian estimation:

- Credible differences between means and standard deviations.
- Complete distributional information on effect size and everything else.
- Non-normality indicated.

NHST t test:

- Outliers invalidate classic test.
- Resampling shows $p > .05$ for difference of means, $p > .05$ for difference of standard deviations.
- Need correction for multiple tests.
- No CI's. (And CI's would have no distributional info and fickle end points linked to fickle p values.)



© John K. Kruschke, Oct. 2012

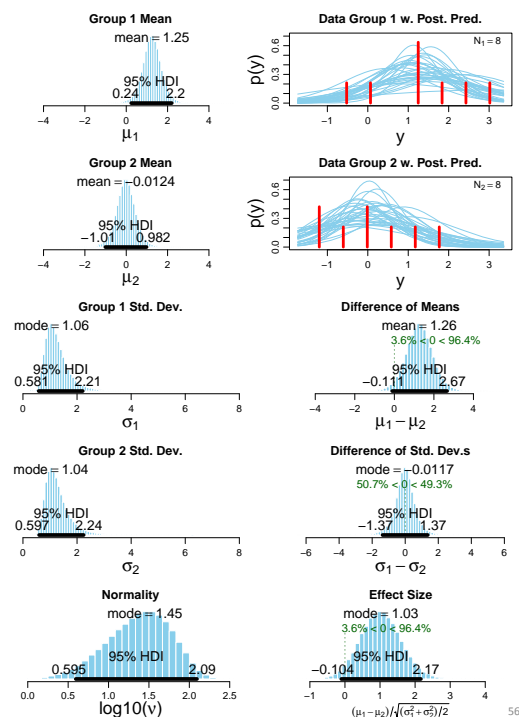
Example with small N

Bayesian estimation:

- **Zero is among credible differences between means** and standard deviations, and for effect size.
- Complete distributional information on effect size and everything else.
- Normality is credible.


NHST t test:

- $t(14) = 2.33$, $p = 0.035$, 95% CI: 0.099, 2.399. ($F(7,7) = 1.00$, $p = .999$, CI on ratio: 0.20, 5.00.)
- Need correction for multiple tests, if intended.
- CI's have no distributional info and fickle end points linked to fickle p values.
- **t test fails to reveal true uncertainty in parameter estimates when simultaneously estimating SD's and normality.**



© John K. Kruschke, Oct. 2012

Region of Practical Equivalence (ROPE)




Consider a landmark value. Values that are equivalent to that landmark for all practical purposes define the ROPE around that value.

For example, the landmark value is 100, and the ROPE is 99 to 101.

© John K. Kruschke, Oct. 2012

57

Region of Practical Equivalence (ROPE)



A parameter value is declared to be not credible, or **rejected**, if its entire ROPE lies outside the 95% HDI of the posterior distribution of that parameter.

A parameter value is declared to be **accepted** for practical purposes if that value's ROPE completely contains the 95% HDI of the posterior of that parameter.

© John K. Kruschke, Oct. 2012

58

Example of accepting null value

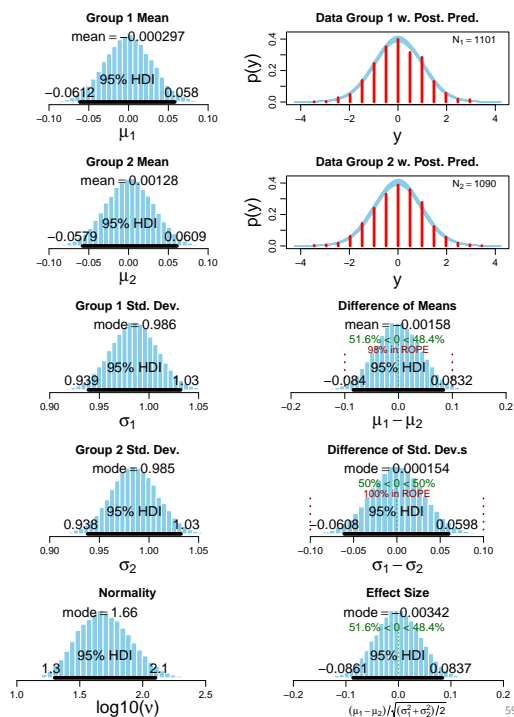
Bayesian estimation:

- 95% HDI for difference on means falls within ROPE; same for SD's (**enlarged in next slide**).
- Complete distributional information on effect size and everything else.
- Normality is credible.

NHST t test:

- p is large for both t and F tests, but NHST cannot accept null hypothesis.
- Need correction for multiple tests, if intended.
- CI's have no distributional info and fickle end points linked to fickle p values, and CI does not indicate probability of parameter value. Hence, **cannot use ROPE method in NHST.**

© John K. Kruschke, Oct. 2012



Example of accepting null value

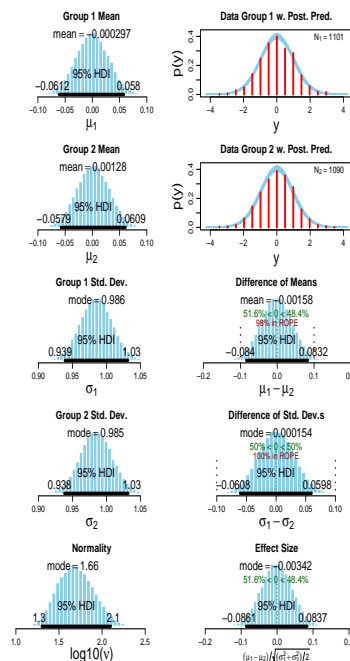
Bayesian estimation:

- 95% HDI for difference on means falls within ROPE; same for SD's.
- Complete distributional information on effect size and everything else.
- Normality is credible.

NHST t test:

- p is large for both t and F tests, but NHST cannot accept null hypothesis.
- Need correction for multiple tests, if intended.
- CI's have no distributional info and fickle end points linked to fickle p values, and CI does not indicate probability of parameter value. Hence, **cannot use ROPE method in NHST.**

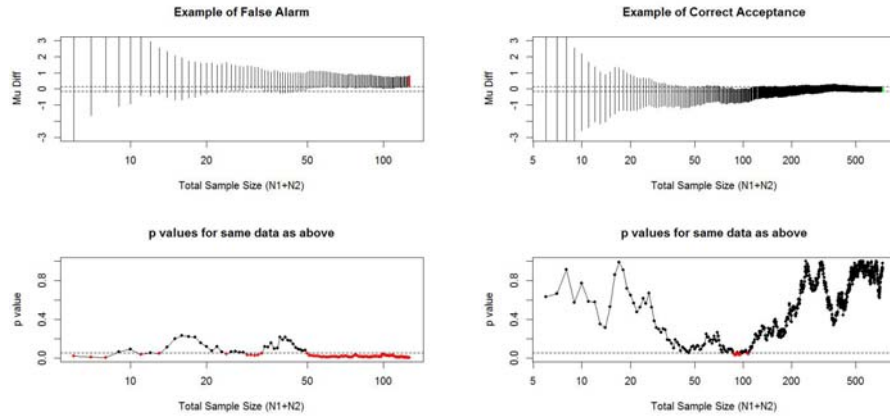
© John K. Kruschke, Oct. 2012



60

Sequential Testing

For simulated data from the null hypothesis:




© John K. Kruschke, Oct. 2012

61

Sequential Testing

For simulated data from the null hypothesis:




© John K. Kruschke, Oct. 2012

62

Many other topics are in the book, e.g.

- ❖ Bayesian hierarchical **ANOVA**, oneway and twoway with interaction contrasts.
- ❖ The **generalized linear model**.
- ❖ Many types of **regression**, including multiple linear regression, logistic regression, ordinal regression.
- ❖ **Log-linear models vs chi-square test**.
- ❖ **Power**: Probability of achieving the goals of research.
- ❖ All preceded by **extensive introductory chapters** covering notions of probability, Bayes' rule, MCMC, model comparison, etc.



© John K. Kruschke, Oct. 2012

63

An example of a t test:

Data:

Group 1: 5.70 5.40 5.75 5.25 4.25 4.74; $M1 = 5.18$

Group 2: 4.55 4.98 4.70 4.78 3.26 3.67; $M2 = 4.32$

$t = 2.33$

Show of hands please:

Who bets that $p < .05$? Who bets that $p > .05$?

© John K. Kruschke, Oct. 2012

64

An example of a t test:

Data:

Group 1: 5.70 5.40 5.75 5.25 4.25 4.74; $M1 = 5.18$

Group 2: 4.55 4.98 4.70 4.78 3.26 3.67; $M2 = 4.32$

$t = 2.33$

Show of hands please:

Who bets that $p < .05$? Who bets that $p > .05$?

You're right!

You're right!

© John K. Kruschke, Oct. 2012

65

Null Hypothesis Significance Testing (NHST)

Consider how we draw conclusions from data:

- Collect data, *carefully insulated from our intentions*.
 - Double blind clinical designs.
 - No datum is influenced by any other datum before or after.
- Compute a summary statistic, e.g., for a difference between groups, the t statistic.
- Compute p value of t . If $p < .05$, declare the result to be "significant."

© John K. Kruschke, Oct. 2012

66

Null Hypothesis Significance Testing (NHST)

Consider how we draw conclusions from data:

- Collect data, *carefully insulated from our intentions*.
 - Double blind clinical d
 - No datum is influence
- Compute a summary s
between groups, the t statistic.

Value of p depends on the
intention of the experimenter!

- Compute p value of t . If $p < .05$, declare the result to be "significant."

© John K. Kruschke, Oct. 2012

67

The road to NHST is paved with good intentions.

The p value is the probability that the actual sample statistic, or a result more extreme, would be obtained from the null hypothesis, *if the **intended** experiment were repeated *ad infinitum*.*

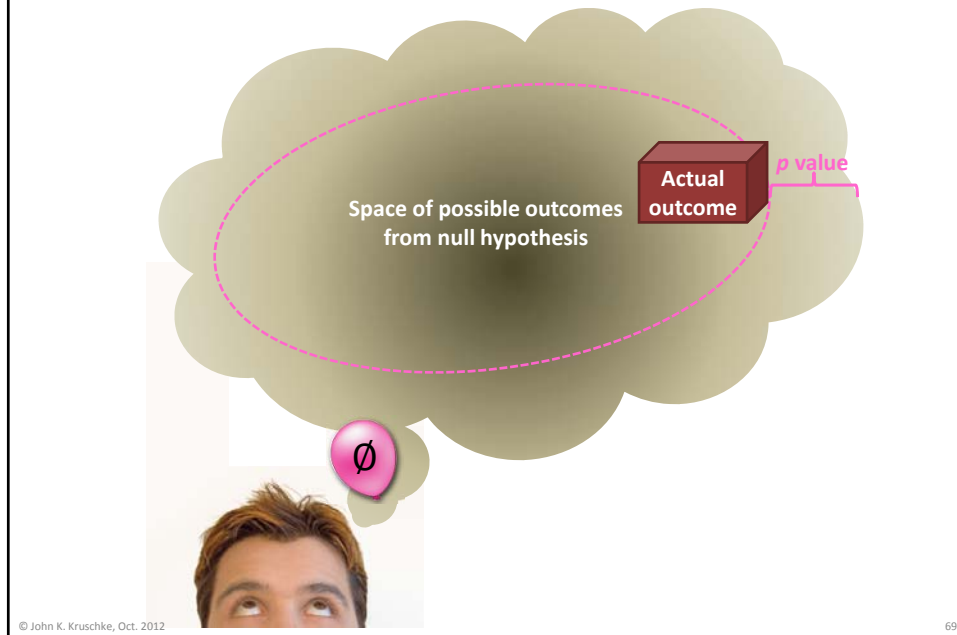
$$p \text{ value} = p(|t_{\text{null}}| > |t_{\text{act}}|)$$

for t_{null} sampled according to
the intended experiment


© John K. Kruschke, Oct. 2012

68


"The" p value...



p value for intention to sample until N

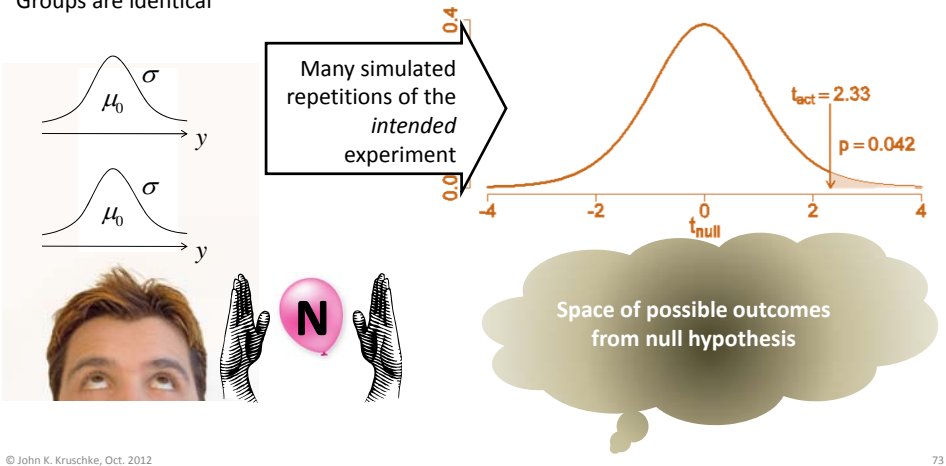


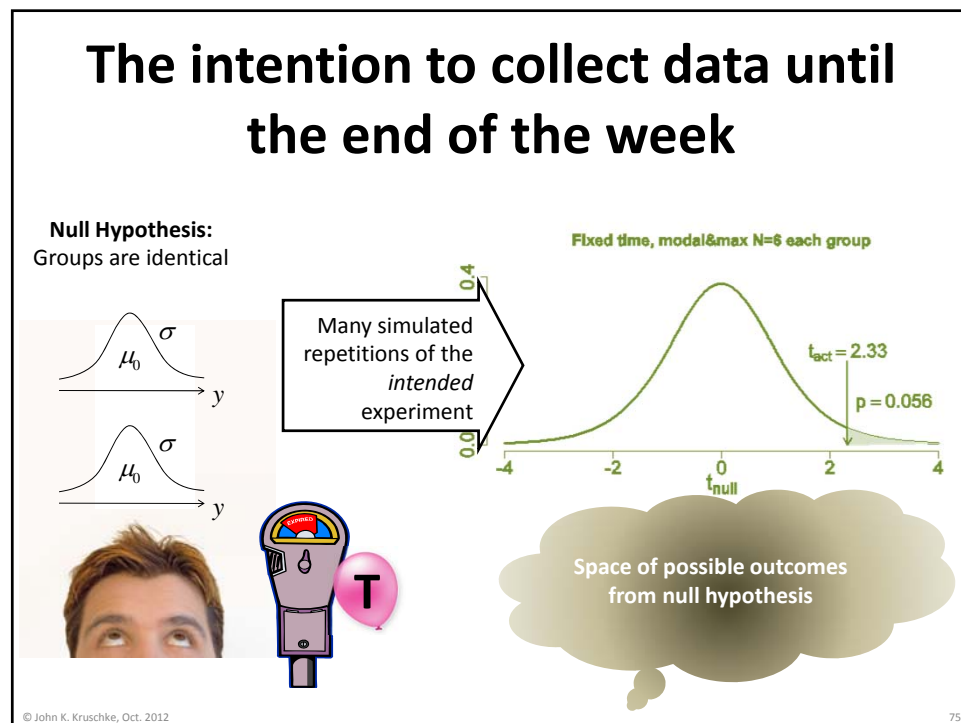

p value for intention to sample until Time



The distribution of t when the intended experiment is repeated many times


Null Hypothesis:
Groups are identical






The intention to collect data until the end of the week

Null Hypothesis:
Groups are identical



Many simulated repetitions of the intended experiment



© John K. Kruschke, Oct. 2012

76

An example of a t test:

Data:

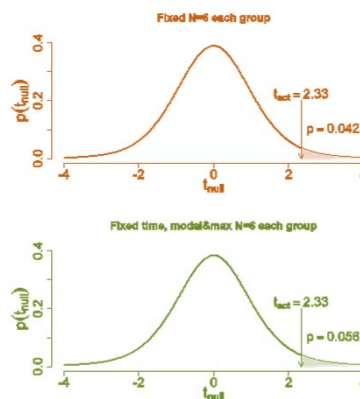
Group 1: 5.70 5.40 5.75 5.25 4.25 4.74; $M1 = 5.18$

Group 2: 4.55 4.98 4.70 4.78 3.26 3.67; $M2 = 4.32$

$t = 2.33$

Can the null hypothesis be rejected? To answer, we must know the intention of the data collector.

- We ask the research assistant who collected the data. The assistant says, "I just collected data for two weeks. It's my job. I happened to get 6 subjects in each group."
- We ask the graduate student who oversaw the assistant. The student says, "I knew we needed 6 subjects per group, so I told the assistant to run for two weeks, because we usually get about 6 subjects per week."
- We ask the lab director, who says, "I told my graduate student to collect 6 subjects per group."
- Therefore, **for the lab director, $t = 2.33$ rejects the null hypothesis (because $p < .05$)**, but for the research assistant who actually collected the data, $t = 2.33$ **fails to reject** the null hypothesis (because $p > .05$).



© John K. Kruschke, Oct. 2012

77

Two labs collect data with same t and N :


Lab A: Collect data until $N=6$ per group.

Data:

Group 1: 5.70 5.40 5.75 5.25 4.25 4.74; $M1 = 5.18$

Group 2: 4.55 4.98 4.70 4.78 3.26 3.67; $M2 = 4.32$

$t = 2.33$



Lab A: Reject the null.


Lab B: Collect data for two weeks.

Data:

Group 1: 5.70 5.40 5.75 5.25 4.25 4.74; $M1 = 5.18$

Group 2: 4.55 4.98 4.70 4.78 3.26 3.67; $M2 = 4.32$

$t = 2.33$



Lab B: Do *not* reject the null.

© John K. Kruschke, Oct. 2012

78

The *real* use of the Neuralyzer:

You *meant* to collect data until $N=12$!

Now *that's* significant!



© John K. Kruschke, Oct. 2012

79

Problem is not solved by “fixing” the intention

- All we need to do is decide in advance exactly what our intention is (or use a Neuralyzer after the fact), and have everybody chant a mantra to keep that intention fixed in their minds while the experiment is being conducted. Right?
- Wrong. The data don't know our intention, and the same data could have been collected under many other intentions.

© John K. Kruschke, Oct. 2012

80

The intention to examine data thoroughly

Many experiments involve multiple groups, and **multiple comparisons** of means.

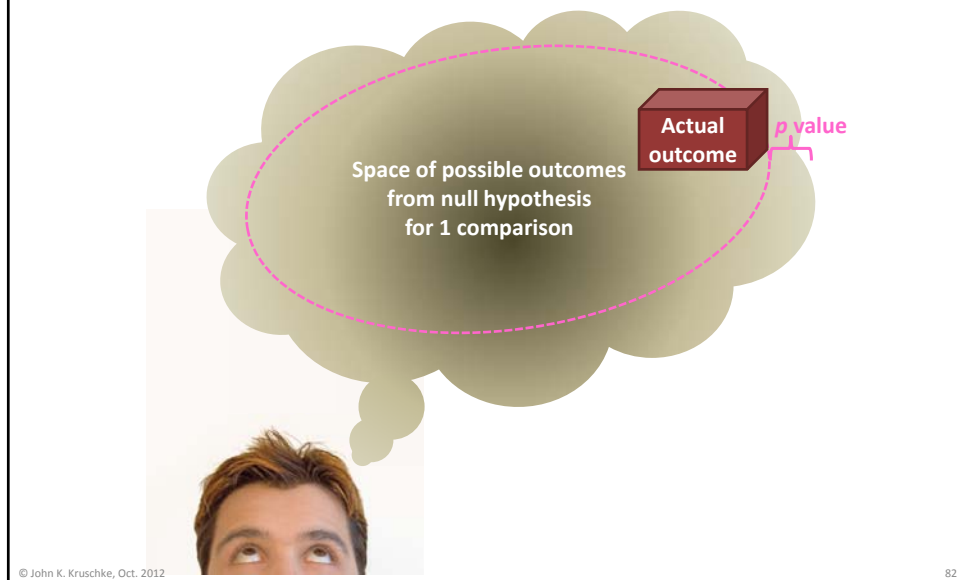
Example: Consider 2 different drugs from chemical family A, 2 different drugs from chemical family B, and a placebo group. Lots of possible comparisons...

Problem: With every test, there is possibility of false alarm! False alarms are bad; therefore, keep the experimentwise false alarm rate down to 5%.

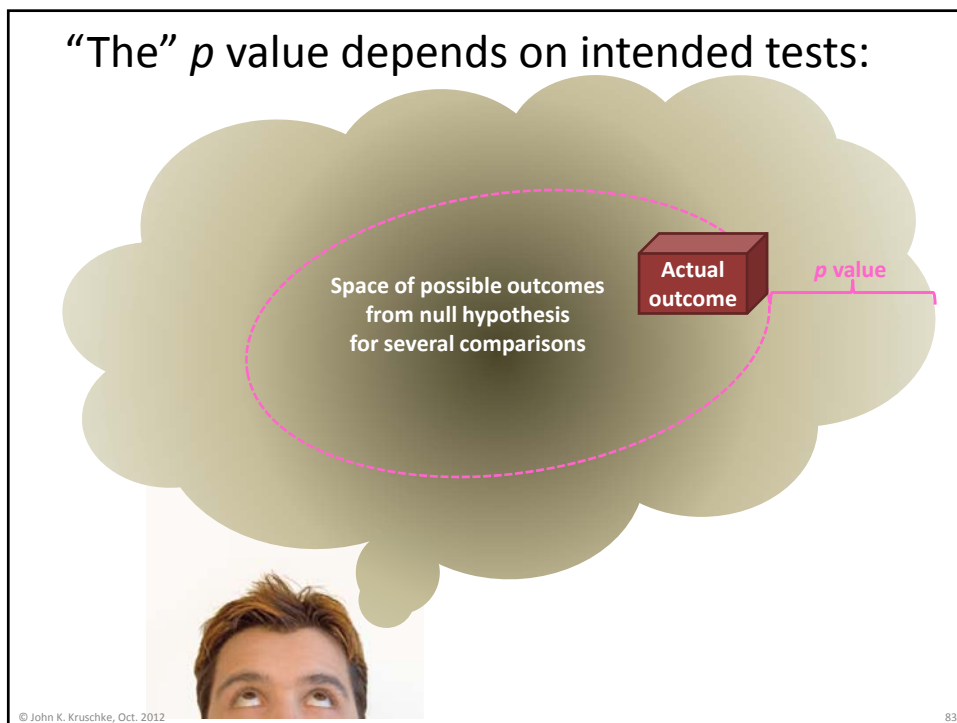
© John K. Kruschke, Oct. 2012

81


“The” p value depends on intended tests:



“The” p value depends on intended tests:



Experimentwise false alarm rate



© John K. Kruschke, Oct. 2012

84

Multiple Corrections for Multiple Comparisons

Begin: **Is goal to identify the best treatment?**

Yes: Use **Hsu's method.**

No: **Contrasts between control group and all other groups?**

Yes: Use **Dunnett's method.**

No: **Testing all pairwise and no complex comparisons (either planned or post hoc) and choosing to test only some pairwise comparisons post hoc?**

Yes: Use **Tukey's method.**

No: **Are all comparisons planned?**

Yes: Use **Scheffe's method.**

No: Is Bonferroni critical value less than Scheffe critical value?

Yes: Use **Bonferroni's method.**

No: Use Scheffe's method (or, prior to collecting the data, reduce the number of contrasts to be tested).

Adapted from Maxwell & Delaney (2004). Designing experiments and analyzing data: A model comparison perspective. Erlbaum.

© John K. Kruschke, Oct. 2012

85

Multiple Corrections for Multiple Comparisons

Begin: *Is goal to identify the best treatment?*

Yes: Use **Hsu's method**.

No: *Contrasts between control group and all other groups?*

Yes: Use **Dunnett's method**.

No: *Testing all pairwise and no complex comparisons (either planned or post hoc) and choosing to test only some pairwise comparisons post hoc?*

Yes: Use **Tukey's method**.

No: *Are all comparisons planned?*

Yes: Use **Scheffe's method**.

No: Is Bonferroni critical value less than Scheffe critical value?

Yes: Use ~~**Bonferroni's method**~~.

! No: Use Scheffe's method (or, prior to collecting the data, reduce the number of contrasts to be tested).

Adapted from Maxwell & Delaney (2004). Designing experiments and analyzing data: A model comparison perspective. Erlbaum.

© John K. Kruschke, Oct. 2012


86

Good intentions make any result *insignificant*

- Consider an experiment with two groups.
- Collect data; compute t test on difference of means. Suppose it yields $p < .05$
- Now, think thoroughly about all the other comparison groups and other experiment groups you should and could meaningfully run.
- Earnestly intend to run them eventually, and to compare your current results with those results.
- *Poof! Your current data are no longer significantly different.*

© John K. Kruschke, Oct. 2012

87



Good intentions make many results *significant*

- Consider an experiment with two groups.
- Collect data; compute t test on difference of means, using df corresponding to actual N . Suppose $p > .05$, but not by much.
- *You had intended to collect a much larger sample size, but you were unexpectedly interrupted.*
- Use the larger intended N for df in the t test.
- *Poof! Your current data are now significantly different!*

? Confidence Intervals provide no confidence ?

Data:

Group 1: 5.70 5.40 5.75 5.25 4.25 4.74; M1 = 5.18

Group 2: 4.55 4.98 4.70 4.78 3.26 3.67; M2 = 4.32

Under assumption of fixed N:

$$\begin{aligned} CI &= (M_1 - M_2) \pm t_{crit} \times se \\ &= (5.18 - 4.32) \pm 2.23 \times 0.370 \\ &= [0.036, 1.68] \end{aligned}$$

which excludes zero.

95% CI constructed with fixed-N t_{crit} will span true difference *less* than 95% of time if data are sampled according to fixed duration.

Under assumption of fixed duration:

$$\begin{aligned} CI &= (M_1 - M_2) \pm t_{crit} \times se \\ &= (5.18 - 4.32) \pm 2.45 \times 0.370 \\ &= [-0.046, 1.77] \end{aligned}$$

which includes zero.

95% CI constructed with fixed-duration t_{crit} will span true difference *more* than 95% of the time if data are sampled according to fixed N.

© John K. Kruschke, Oct. 2012

90

? Confidence Intervals provide no confidence ?

General definition of CI:

95% CI is the range of parameter values (e.g., $\mu_1 - \mu_2$) that would not be rejected by $p < .05$

Hence, *the 95% CI is as ill-defined as the p value.*

We see this dramatically in confidence intervals corrected for multiple comparisons.

© John K. Kruschke, Oct. 2012

91

? **Confidence Intervals** ? provide no confidence

Confidence intervals provide no distributional information:

We have no idea whether a point at the limit of the confidence interval is any less credible than a point in the middle of the interval.

Implies

vast range for predictions of new data, and "virtually unknowable" power.

© John K. Kruschke, Oct. 2012

92

NHST autopsy

- p values are ill-defined: depend on sampling intentions of data collector. Any set of data has many different p values.
- Confidence intervals are as ill-defined as p values because they are defined in terms of p values.
- Confidence intervals carry no distributional information.

© John K. Kruschke, Oct. 2012

93

Bayesian Estimation or NHST?

When Bayesian estimation and NHST *agree*, which should be used?

Bayesian estimation gives the most complete and informative answer. Answer from NHST is not informative and is fickle.

When Bayesian estimation and NHST *disagree*, which should be used?

Bayesian estimation gives the most complete and informative answer. Answer from NHST is not informative and is fickle.

© John K. Kruschke, Oct. 2012

94

Conclusion

- *p values are not well defined, nor are the limits of confidence intervals, and confidence intervals have no distributional info.*
- Bayesian data analysis is the most complete and normatively correct way to estimate parameters in any model, for all your data.
- Bayesian data analysis is taking hold in 21st century science, from astronomy to zoology. *Don't be left behind.*
- And, for more info, ...

© John K. Kruschke, Oct. 2012

95

The blog: <http://doingbayesiandataanalysis.blogspot.com/>



Kruschke, J. K. (2012). Bayesian estimation supersedes the t test. *Journal of Experimental Psychology: General*.




Kruschke, J. K. (2011). Bayesian assessment of null values via parameter estimation and model comparison. *Perspectives on Psychological Science*, 6(3), 299-312.



Kruschke, J. K. (2010). What to believe: Bayesian methods for data analysis. *Trends in Cognitive Sciences*, 14(7), 293-300.



Kruschke, J. K. (2010). Bayesian data analysis. *Wiley Interdisciplinary Reviews: Cognitive Science*, 1(5), 658-676.




Kruschke, J. K. (2011). *Doing Bayesian Data Analysis: A Tutorial with R and BUGS*. Academic Press / Elsevier.

© John K. Kruschke, Oct. 2012

96

Program and manuscript at
<http://www.indiana.edu/~kruschke/BEST/>



© John K. Kruschke, Oct. 2012

97

Priors are not capricious

1. Priors are explicitly specified and must be acceptable to a skeptical scientific audience.
2. Typically, priors are set to be noncommittal and have very little influence on the posterior.
3. Priors can be informed by well-established data and theory, thereby giving inferential leverage to small samples.
4. When there is disagreement about the prior, then the influence of the prior on the posterior can be, and is, directly investigated. Different theoretically-informed priors can be checked.
5. Not using priors can be a serious blunder! E.g., drug/disease testing without incorporating prior knowledge of base rates.

© John K. Kruschke, Oct. 2012

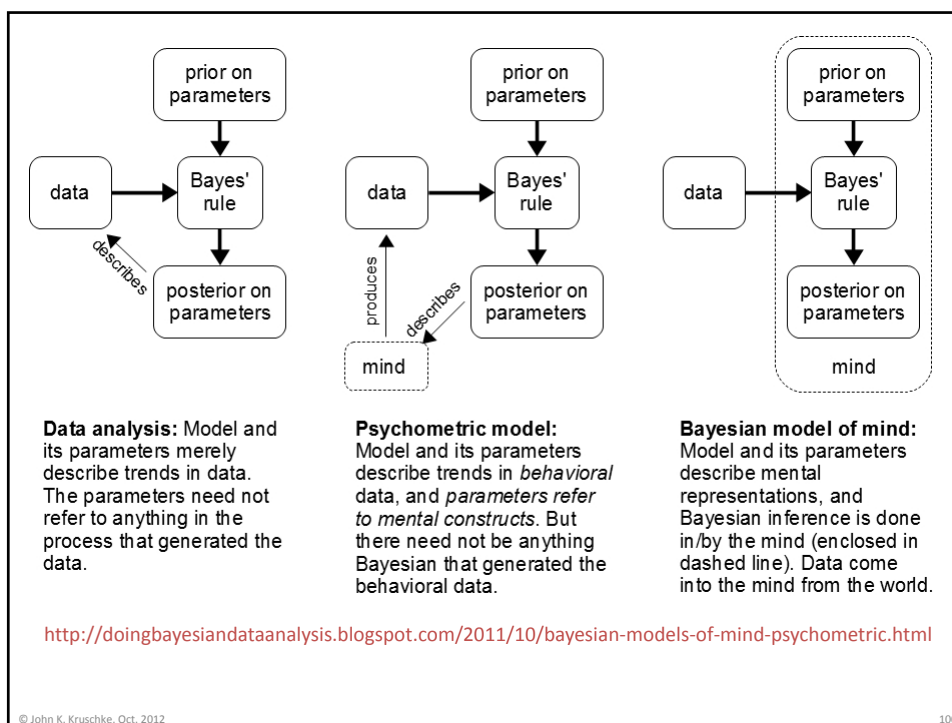
98

Prior credibility is *not* intentions

Bayesian Prior	NHST Intention (e.g., stopping rule, number of comparisons)
Explicit and supported by previous data.	Unknowable
Should influence interpretation of data.	Should <i>not</i> influence interpretation of data

© John K. Kruschke, Oct. 2012

99




Bayesian estimation or Bayesian model comparison?

Bayesian estimation is also better than the “Bayesian t test,” which uses the “Bayes factor” from Bayesian model comparison...



Kruschke, J. K. (2011). Bayesian assessment of null values via parameter estimation and model comparison. *Perspectives on Psychological Science*, 6(3), 299-312.

Chapter 12 of Kruschke, J. K. (2011). *Doing Bayesian Data Analysis: A Tutorial with R and BUGS*. Academic Press / Elsevier.



Kruschke, J. K. (in press). Bayesian estimation supersedes the t test. *Journal of Experimental Psychology: General*. Appendix D.